

Is Your AI Ready to Survive an Audit?

The 5-Minute AI-Readiness Check + Worksheet

Your team is already using AI. The question is whether you could **prove what it did** if a client, a regulator, or opposing counsel asked. Answer each question below — score **1 point** for every “No / Not sure.”

PART 1 — THE CHECK

1. **The perimeter.** Can you name exactly where your data goes when your team uses an AI tool? ☐ Yes ☐ No/Not sure
2. **Shadow AI.** Could a staff member paste a confidential document into a free AI tool today without anyone knowing? ☐ No, covered ☐ Yes/Not sure
3. **Vendor agreement.** For any AI tool touching sensitive data, has the vendor signed a BAA / DPA? ☐ Yes ☐ No/Not sure
4. **Audit trail.** Could you produce a log of every question your AI was asked about a matter and what it produced? ☐ Yes ☐ No/Not sure
5. **Grounding.** Does your AI answer from your real documents — or sometimes make things up? ☐ Grounded ☐ Makes things up/Not sure
6. **Human in the loop.** Is there a human reviewing AI output before it reaches a client or drives a decision? ☐ Always ☐ Not consistently
7. **The policy.** Do you have a written AI-use policy your team knows and follows? ☐ Yes ☐ No/Not sure
8. **The rulebook.** Do you know which rules govern your AI use (Rule 1.6, HIPAA, GDPR, FINRA)? ☐ Yes ☐ No/Not sure
9. **Access control.** Is access to AI systems and the data they reach limited by role? ☐ Yes ☐ No/Not sure
10. **Ownership of output.** Do you know who owns, and who can see, what your AI produces? ☐ Yes ☐ No/Not sure

YOUR SCORE

0–2 — On Track (but verify). Ahead of most. The risk is the exposure you can’t see. A short review confirms it and catches blind spots.

3–6 — At Risk. Real, specific exposure. One pasted document or client questionnaire from a problem. The most common, most fixable range.

7–10 — Exposed. Likely already out of compliance somewhere, and you couldn’t prove otherwise. Put a defensible version in place fast.

Whatever your score, the next step is the same: knowing you have gaps and knowing *which 2–3 to fix first* are very different. That’s the AI-Readiness Review.

PART 2 — THE WORKSHEET

Fill this in and you've done the first hour of the work — it's exactly what we build out together in a Readiness Review.

A. Tool inventory

Tool	Who uses it	Data it could touch	Vendor agreement?

B. Sensitive-data map

What protected / regulated information exists? (client files, patient records, financials, privileged material, PII, IP)

C. The 5-component gap check

Rate each pillar: R = missing Y = partial G = solid

Component	What it means	R Y G
Data perimeter	Sensitive data stays inside your control	■ ■ ■
Grounded retrieval	Answers from your real documents, not guesses	■ ■ ■
Audit trail	Every query, source, and output is logged	■ ■ ■
Human in the loop	A person reviews consequential output	■ ■ ■
Governance	Policy, access control, and ownership are defined	■ ■ ■

D. Your top 3 exposures

E. Where AI would actually help (2 workflows)

Now make it a plan.

The AI-Readiness Review turns this into a prioritized roadmap and a scoped build proposal — a 1–2 week engagement, fee credited toward your build. The first step is a free intro call.

Book a free intro call: calendly.com/kylecedrichanna/ai-readiness-intro · kyle@kylecedrichanna.com